

Question1

(a) # of tuples: $80000 * (1 / 5) * ((600 - 100) / 10000) = 800$

Tuples / leaf: $80000 / 500 = 160$

Cost of index search = $2 + 800 / 160 (+ 1) = 7 (+ 1)$

(b) # of pages: $800 * (5 * 8) / 6000 = 6$

Cost: $\text{ceil}(6 / 23) * 4000 = 4000$

(c) # of tuples: $800 * 600000 * (1 / 5000) = 96000$

of pages: $(96000) * (8 * 6) / 6000 = 768$ pages

$768 / 100 = 8$ groups

The whole process: r/w/r

We do not need to have any cost on read, $\text{cost}(w) + \text{cost}(r) = 768 + 768 = 1536$ pages

(d) Total cost: $7 (+1) + 4000 + 1536 = 5543 (+1)$ pages

Question2:

1. (R join S) join T

First cost (R join S): $200 + (200 / 200) * 1200 = 1400$

For cost to join T: $(8000 / 200) * 4000 = 160000$

Total cost: $160000 + 1400 = 161400$ pages

2. (R join T) join S

First cost (R join T): $200 + (200 / 200) * 4000 = 4200$

For cost to join S: $(2000 / 200) * 1200 = 12000$

Total cost: $12000 + 4200 = 16200$ pages

Question3:

(a) List all conflicts:


$r1(x) w3(x)$

$r2(z) w3(z)$

$w2(w) r1(w)$

$w2(z) r3(z)$

$w2(z) w3(z)$



(b) From the graph above, it is serializable. Since we can go from T2 , T1 , T3 with no cycle.

(c) Yes, it is possible for 2PL.

$r_2(z) \ w_2(w) \ w_2(z) \ r_1(x) \ r_1(y) \ r_1(w) \ w_1(y) \ r_3(z) \ w_3(x) \ w_3(z)$

Simple go from T2 T1 T3, once one finished and do not need any lock, we give the lock to another transaction.

Question4

Query1:

(a)
 SELECT DISTINCT
 a. name
 FROM
 artists a
 , songs s
 , spotify sp
 WHERE
 a.id = s.artistid
 and s.id = sp.songid
 and sp.streams >= 10000000
 ORDER BY
 a. name ASC;

(b)
 Unique (cost=564.30..565.11 rows=163 width=11)
 -> Sort (cost=564.30..564.71 rows=163 width=11)
 Sort Key: a.name
 -> Hash Join (cost=43.54..558.31 rows=163 width=11)
 Hash Cond: (s.artistid = a.id)
 -> Hash Join (cost=25.23..539.57 rows=163 width=8)
 Hash Cond: (sp.songid = s.id)

-> Seq Scan on spotify sp (cost=0.00..513.91 rows=163 width=8)
 Filter: (streams >= 100000000)
 -> Hash (cost=16.77..16.77 rows=677 width=16)
 -> Seq Scan on songs s (cost=0.00..16.77 rows=677 width=16)
 -> Hash (cost=10.36..10.36 rows=636 width=19)
 -> Seq Scan on artists a (cost=0.00..10.36 rows=636 width=19)

(c)

create index query6 on spotify(streams)

(d)

Unique (cost=236.40..237.22 rows=163 width=11)
 -> Sort (cost=236.40..236.81 rows=163 width=11)
 Sort Key: a.name
 -> Hash Join (cost=49.09..230.42 rows=163 width=11)
 Hash Cond: (s.artistid = a.id)
 -> Hash Join (cost=30.78..211.68 rows=163 width=8)
 Hash Cond: (sp.songid = s.id)
 -> Bitmap Heap Scan on spotify sp (cost=5.55..186.02 rows=163 width=8)
 Recheck Cond: (streams >= 100000000)
 -> Bitmap Index Scan on query6 (cost=0.00..5.51 rows=163 width=0)
 Index Cond: (streams >= 100000000)
 -> Hash (cost=16.77..16.77 rows=677 width=16)
 -> Seq Scan on songs s (cost=0.00..16.77 rows=677 width=16)
 -> Hash (cost=10.36..10.36 rows=636 width=19)
 -> Seq Scan on artists a (cost=0.00..10.36 rows=636 width=19)

Query2:

(a) select
 s.name as songname
 , a.name as artistname
 , count(distinct p.id) as numplayed
 from
 songs s
 join billboard b on b.songid = s.id
 join artists a on a.id = s.artistid
 left join playedonradio p on p.songid = s.id
 where

```

    s.danceability >= 0.9
    and b.rank <= 10
group by
    s.id
    , s.name
    , a.name
order by
    numplayed desc
    , songname asc
    , artistname asc
;

```

(b)

```

Sort (cost=1291.57..1294.15 rows=1033 width=45)
Sort Key: (count(DISTINCT p.id)) DESC, s.name, a.name
-> GroupAggregate (cost=1162.07..1239.85 rows=1033 width=45)
Group Key: s.id, a.name
-> Incremental Sort (cost=1162.07..1221.77 rows=1033 width=41)
Sort Key: s.id, a.name
Presorted Key: s.id
-> Merge Join (cost=1161.90..1181.71 rows=1033 width=41)
Merge Cond: (s.id = b.songid)
-> Sort (cost=995.89..997.27 rows=551 width=41)
Sort Key: s.id
-> Hash Join (cost=36.87..970.80 rows=551 width=41)
Hash Cond: (s.artistid = a.id)
-> Hash Right Join (cost=18.56..951.03 rows=551 width=38)
Hash Cond: (p.songid = s.id)
-> Seq Scan on playedonradio p (cost=0.00..809.30 rows=46630
width=12)
-> Hash (cost=18.46..18.46 rows=8 width=34)
-> Seq Scan on songs s (cost=0.00..18.46 rows=8 width=34)
Filter: (danceability >= '0.9'::double precision)
-> Hash (cost=10.36..10.36 rows=636 width=19)
-> Seq Scan on artists a (cost=0.00..10.36 rows=636 width=19)
-> Sort (cost=166.01..169.27 rows=1305 width=8)
Sort Key: b.songid
-> Seq Scan on billboard b (cost=0.00..98.47 rows=1305 width=8)
Filter: (rank <= 10)

```

(c) create index query1 on songs(danceability);
create index query1rank on billboard(rank);

(d)

Sort (cost=1260.83..1263.42 rows=1033 width=45)

Sort Key: (count(DISTINCT p.id)) DESC, s.name, a.name

-> GroupAggregate (cost=1131.34..1209.12 rows=1033 width=45)

Group Key: s.id, a.name

-> Incremental Sort (cost=1131.34..1191.04 rows=1033 width=41)

Sort Key: s.id, a.name

Presorted Key: s.id

-> Merge Join (cost=1131.16..1150.97 rows=1033 width=41)

Merge Cond: (s.id = b.songid)

-> Sort (cost=991.92..993.30 rows=551 width=41)

Sort Key: s.id

-> Hash Join (cost=32.90..966.83 rows=551 width=41)

Hash Cond: (s.artistid = a.id)

-> Hash Right Join (cost=14.59..947.07 rows=551 width=38)

Hash Cond: (p.songid = s.id)

-> Seq Scan on playedonradio p (cost=0.00..809.30 rows=46630 width=12)

-> Hash (cost=14.49..14.49 rows=8 width=34)

-> Bitmap Heap Scan on songs s (cost=4.34..14.49 rows=8 width=34)

Recheck Cond: (danceability >= '0.9'::double precision)

-> Bitmap Index Scan on query1 (cost=0.00..4.33 rows=8 width=0)

Index Cond: (danceability >= '0.9'::double precision)

-> Hash (cost=10.36..10.36 rows=636 width=19)

-> Seq Scan on artists a (cost=0.00..10.36 rows=636 width=19)

-> Sort (cost=139.24..142.50 rows=1305 width=8)

Sort Key: b.songid

-> Bitmap Heap Scan on billboard b (cost=18.40..71.71 rows=1305 width=8)

Recheck Cond: (rank <= 10)

-> Bitmap Index Scan on query1rank (cost=0.00..18.07 rows=1305 width=0)

Index Cond: (rank <= 10)